

CS 484 - MP1

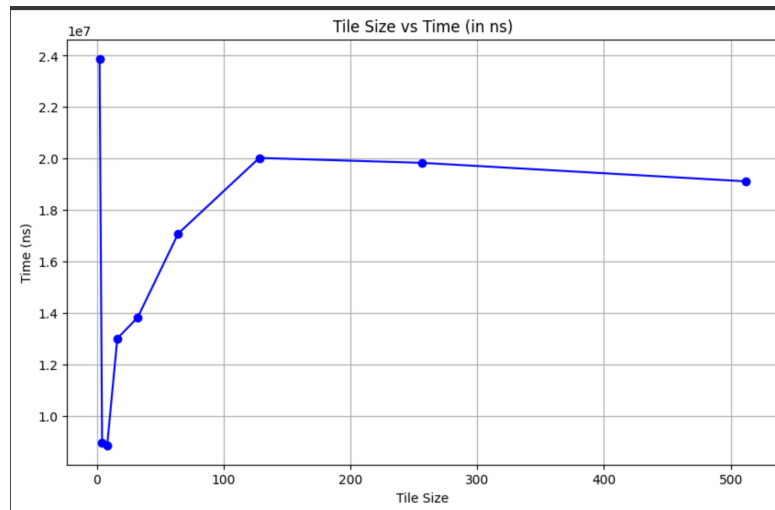
NetId - Harsha4

Date - 02/17/24

Link to Google Colab - [CS484 Anaylsis.ipynb](#)

1.1

1.

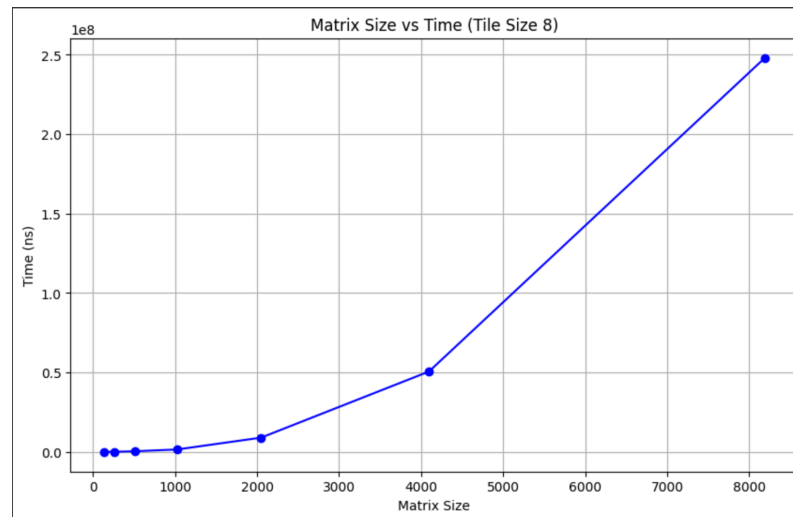


From the `cache_properties.txt` file, we see the following information

LEVEL1_ICACHE_LINESIZE	64
LEVEL1_ICACHE_ASSOC	8

This signifies that each cache line is of size 64 bytes. Since the array is of doubles, there can be 8 elements in one cache line. The set associativity of this cache is 8 which means that there are 8 cache lines in total. As a result, using tile size of 8 will provide maximum ability to exploit the spatial cache locality and hence take the least amount of time.

2.



The above graph shows the time taken by varying sizes of N to complete execution. Specifically, using a tile size of 8 fits well with the machine's specifications since the cache line is 64 bytes long and can fit 8 doubles. Having 8 way set associativity also means it can fit 8 cache lines in the L-1 cache. As a result, when the tile size is 8, the spatial locality is extensively used and hence helps in more cache hits making 8 an optimal choice.

Conclusion

Using a tile of size 8 improves performance by over 4x when the array is large enough.

1.2

Benchmark	Time	CPU
BM_Transpose_Naive/matrix_size:512/process_time/real_time	888326 ns	888252 ns
BM_Transpose_Tiled/matrix_size:512/tile_size:8/process_time/real_time	214977 ns	214961 ns
BM_Multiply_Naive/matrix_size:512/process_time/real_time	444028039 ns	443994178 ns
BM_Multiply_Tiled/matrix_size:512/tile_size:8/process_time/real_time	100699554 ns	100691922 ns
BM_Multiply_Transposed/matrix_size:512/process_time/real_time	87305868 ns	87296634 ns

From: gbench_results.txt

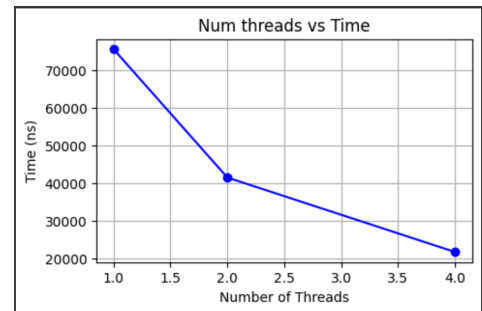
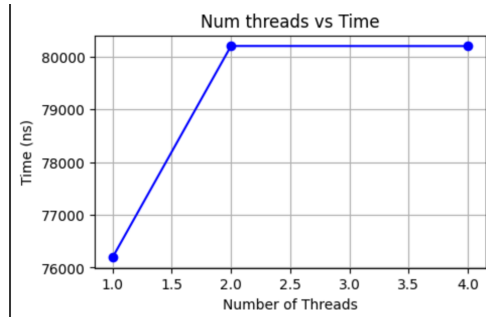
```
≡ multiply_results.txt ×
mp1 > writeup > ≡ multiply_results.txt
1 Average execution time of multiply_basic(): 1.007680
2 Average execution time of multiply_tiled(): 0.131155
3 Average execution time of multiply_transposed(): 0.088673
4
```

From: multiply_results.txt

Based on the tests and benchmark tools used above, it is evidently clear that **Multiply_Transposed** works better than **Multiply_Tiled**. This can be explained as follows - Transposing one of the matrices can enhance memory access patterns by improving spatial locality. When accessing elements from the transposed matrix, consecutive memory accesses are more likely to be adjacent in memory, which can reduce cache misses and improve performance.

1.3

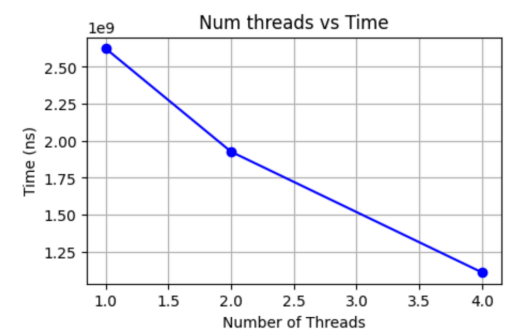
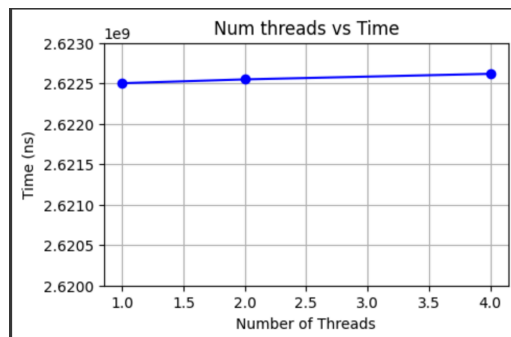
1.



Function: `squares_parallel`

Without OpenMP

With OpenMP



Function: `primes_parallel`

Without OpenMP

With OpenMP

In both functions, Using 4 threads reduces the time drastically.

For the first function, the result of using 4 threads increased speedup by 4X and for the second function, the result of using 4 threads increased speedup by 2X. Hence `#pragma omp parallel for` increases speedup by 4X and 2X respectively.